

BENJAMIN OFURHIE

Lisbon, Portugal | OfurhieB@yahoo.com | Phone: +35192005674 | [LinkedIn](#) | [Portfolio](#)

Data Engineer with experience building ETL pipelines, cloud data workflows, and analytics solutions using Python, SQL, Airflow, PySpark, GCP, and AWS. Currently developing production data systems at Dealroom and delivering end-to-end data engineering projects covering orchestration, warehousing, data quality, and cloud analytics.

Technical Skills

- Programming Languages: Python, SQL
- Cloud & Platforms: AWS, Google Cloud Platform, Databricks
- Databases: PostgreSQL, MySQL, BigQuery, MongoDB
- Tools: dbt (Data Build Tool), Apache Airflow, Docker, GitHub Actions (CI/CD), Git, PySpark
- Data Quality: Data Validation, Pipeline Monitoring, Error Handling
- Visualization: Power BI, Looker Studio, Tableau

Education & Certification

- Google Cloud Certified Professional Data Engineer | Google Cloud
- Apache Airflow 3 Fundamentals Certification | Astronomer
- Data Engineering Professional Certificate | IBM
- AWS Certified Cloud Practitioner | Amazon Web Services (AWS)

Experience

Data Operations Specialist | Dealroom.co

Lisbon, Portugal | 2024 -Present

- Built Python scraping scripts to extract Y Combinator batch data, automating enrichment of 20–30 new companies per run, executed twice weekly against the main database.
- Designed and maintained Python-based ETL processes to extract, clean, and transform API data feeding team-wide dashboards and KPI reporting.
- Scheduled and automated data workflows using GCP Cloud Scheduler to ensure reliable and timely data delivery.
- Developed custom automation tools utilizing OpenClaw to streamline data research workflows and handle automated quality checks.
- Wrote and optimized SQL queries across large-scale BigQuery tables to support reporting and metric definition.

Data Analyst | Meri Skill

Remote | 2022–2023

- Designed and automated ETL workflows using Python (Pandas) and SQL to process, clean, and structure large datasets for operational performance analysis.
- Built and maintained data pipelines using Power Query and SQL to standardize raw data and ensure reliable, timely delivery of recurring business reports.
- Developed Python scripts to automate data ingestion and transformation, reducing manual processing time across operational datasets.

Data Engineering Projects

YouTube ETL Pipeline | Python, PostgreSQL, Docker, Airflow Soda, GitHub Actions [Github Repo](#)

- Built an end-to-end ETL pipeline extracting 986 videos from the MrBeast channel via the YouTube API, with daily incremental loads using unique video IDs as deduplication constraints.
- Designed staging and core layers in PostgreSQL, implementing custom SQL transformations and relational data modeling practices.
- Developed Python scripts to handle automated ingestion, schema validation, flattening of JSON data, and data cleaning.
- Implemented automated data quality checks using Soda alongside Python unit tests to validate pipeline logic and data integrity.
- Containerized the application with Docker and built a GitHub Actions CI/CD workflow to automate testing and production deployment.

Uber ETL Pipeline | Python, Airflow, BigQuery, GCP, Docker [Github Repo](#)

- Designed a scalable ETL pipeline processing ~4.5M rows of Uber ride data, loading securely into Google Cloud Storage for downstream analytics.
- Orchestrated automated daily workflows using Apache Airflow, implementing custom DAGs with retry logic and error handling.
- Developed Python (Pandas) validation routines to clean raw ingestion batches, handle null constraints, and standardize data schemas.
- Modeled and structured transformed datasets into BigQuery analytics layers, creating optimized tables for performance reporting.
- Containerized the entire Airflow environment using Docker, ensuring fully reproducible deployment and execution across cloud environments.

Data Warehouse/Modeling Project | MySQL, SQL Server, Medallion Architecture [Github Repo](#)

- Engineered an end-to-end SQL data warehouse ingesting data from 2 source systems (CRM and ERP), delivering a star schema with 3 analytics-ready tables (fact_sales, dim_products, dim_customers) across Bronze/Silver/Gold layers.
- Designed a star schema model with conformed dimensions and fact tables to optimize analytical query performance..
- Developed complex SQL transformations using CTEs, window functions, joins, and conditional logic for data cleaning and enrichment.
- Implemented data quality controls including deduplication, validation rules, and schema standardization.

Sales ETL Pipeline | Python, PySpark, Airflow, PostgreSQL, Docker [Github Repo](#)

- Built an end-to-end ETL pipeline processing 32,718 sales records using PySpark, orchestrated with Apache Airflow on a fully containerised Docker environment.
- Engineered a Parquet-based staging architecture decoupling extract and transform layers for independent re-execution and fault tolerance.
- Applied PySpark transformations including explicit type casting, CamelCase-to-snake_case column standardisation, and derived date field enrichment.
- Implemented an idempotent UPSERT loading strategy using ON CONFLICT to ensure safe re-triggers without duplicate records.